# 1. Getting started

## 1.1 Introduction

FLAGR is a high performance, modular, open source library for rank aggregation problems. It implements baseline and recent state-of-the-art aggregation algorithms that accept ranked preference lists and generate a single consensus list of elements.

The core project is developed in C++. The source code is available on GitHub and can be compiled as a standard application, or as a shared library. In the second case, the library file can be linked or loaded by other programs in other languages. PyFLAGR is an example of such application. In brief, FLAGR:

- employs efficient data structures and algorithms that ensure high performance,
- is cross-platform supporting Windows and Linux. An extension that supports MacOS users is currently under construction,
- is modular, allowing third-party programmers to easily implement their methods within the core library,
- is open-source.

The current version of FLAGR is 1.0.7. It includes implementations of the following algorithms:

- CombSUM linear combination with 5 different score/rank normalization techniques; namely: Rank, Borda, Simple Borda, Score, and Z-Score normalization (Renda, et al., 2003).
- CombMNZ linear combination with 5 different score/rank normalization techniques; namely: Rank, Borda, Simple Borda, Score, and Z-Score normalization (Renda, et al., 2003).
- Borda Count (equivalent to CombSUM with Borda normalization, (Renda, et al., 2003)).
- Condorcet Winners.
- Copeland Winners.
- Outranking Approach of Farah & Vanderpooten, 2007.
- Distance-based iterative unsupervised algorithm of Akritidis et al., 2022 (all the above methods can be used as the starting non-weighted aggregator).
- Robust Rank Aggregation in two variants: the first one employs the Stuart/Ares method for p-value correction, whereas the other one does not.
- Kemeny optimal aggregation (brute force implementation, not applicable to large, or many input preference lists).
- Markov Chains (MC) methods of Dwork et al., 2001 and DeConde et al., 2006.
- Weighted agglomerative aggregation method of Chatterjee et al., 2018.
- Preference relation unsupervised algorithm of Desarkar et al., 2016.

These methods are also supported by PyFLAGR through a set of modules and classes.

# 1.2 Download

FLAGR is an open-source project, licensed under the Apache License, version 2.

The library can be downloaded from its GitHub repository at https://github.com/lakritidis/FLAGR. The repository contains:

- The core C++ components (`/src` directory) and rank aggregation algorithm implementations (`/src/ram` directory).

- The dynamic library references (`cflagr.cpp` and `dllflagr.cpp` files), which allow the compilation of FLAGR as a shared or a dynamic link library.

- A precompiled shared library (`pyflagr/pyflagr/flagr.so`, Linux only) and a dynamic link library (`pyflagr/pyflagr/flagr.dll`, Windows only) with its dependencies.

- The PyFLAGR library, which allows the usage of FLAGR in Python applications (`/pyflagr` directory).

- A Jupyter notebook that demonstrates how the PyFLAGR modules can be imported and used in Python applications (`README.ipynb`).

# 1.3 Compilation

FLAGR can be compiled as a standard console application, or as a shared/dynamic link library.

## Building FLAGR as a standard console application

In Linux platforms, GCC can be used to compile `cflagr.cpp` as follows:

```
g++ -Wall -std=gnu++11 -O3 -std=gnu++11 -Isrc/ -c /path/to/cflagr/cflagr.cpp -o /path/to/obj/flagr.o

g++  -o /path/to/out/FLAGR obj/Release/cflagr.o  -O3 -s
```

For Windows-based systems, MingW can be used in a similar way:

```
g++ -Wall -std=gnu++11 -O3 -std=gnu++11 -Isrc\ -c /path/to/cflagr/cflagr.cpp -o /path/to/obj/flagr.o

g++ -o /path/to/out/FLAGR.exe /path/to/obj/cflagr.o  -O3 -s
```

## Building FLAGR as a shared library or a dynamic link library

At first, please notice that the GitHub repository of FLAGR already contains fully-functional pre-built shared and dynamic link libraries that have been created and tested with the GCC compiler and the Valgrind memory profiler. Both of them are located in the `pyflagr/pyflagr` directory and can be used immediately in third party applications (do not forget to also include the required dependencies).

In cases where a manual build is required (e.g., you have implemented your own rank aggregation or evaluation method), GCC can be used with the `-shared` flag as follows:

```
g++ -O3 -Wall -Werror -shared -std=c++11 -fPIC /path/to/cflagr/cflagr.cpp -o
/path/to/so/flagr.so
```

This command will generate the necessary `.so` Linux shared library.

For Windows-based systems with the GCC compiler, FLAGR can be built as a Dynamic Link Library by invoking the following system commands:

```
g++ -O3 -c -o flagr.o /path/to/cflagr/dllflagr.cpp
g++ -O3 -o /path/to/dll/flagr.dll -s -shared flagr.o -Wl,--subsystem,windows
```

This command will generate the necessary `.dll` library. An automated procedure for compiling FLAGR as a DLL can be executed from the `makedll.bat` batch file included in the GitHub repository. Please make sure that you have specified the correct file locations in `makedll.bat` before executing it.

# 1.4 Dynamic library references (Exposed C functions)

FLAGR exposes a set of C functions through an `extern "C"` statement, allowing their linkage from other C programs. These are the dynamic library references to which a client C code can link to. At first, FLAGR must be compiled as a shared/dynamic link library. Then, the client C program can link to that shared library and include a typical C header file that contains just the declaration of these functions. The called function is able to access the C++ FLAGR core through the Execution Driver. The following diagram depicts this scenario:



The FLAGR architecture and the possibility of building it as a shared library allows its usage not only in third party C programs, but also in programs written in other languages. PyFLAGR is an example of such case. PyFLAGR is a Python library that enables the execution of the algorithm C++ implementations of FLAGR from standard Python programs. The Python program must be able to successfully import the PyFLAGR modules.

The exposed C functions of FLAGR exist in two files: `cflagr.cpp` and `dllflagr.cpp`. These files are almost identical: they contain the same functions with exactly the same arguments and bodies. What changes is the usage of several special keywords in the function declarations of `dllflagr.cpp` that enable the building of FLAGR as a DLL for Windows-based systems. These functions are:

- `void Linear()`: It executes one of the supported linear combination methods (CombSUM and CombMNZ, each one with 5 variants).

- `void Condorcet()`: It executes the Condorcet Winners method.

- **void `Copeland()`**: It executes the Copeland Winners method.

- **void `OutrankingApproach()`**: It executes the Outranking Approach of Farah & Vanderpooten, 2007.

- **void `Kemeny()`**: It executes the optimal Kemeny optimal aggregation algorithm (brute force implementation).

- **void `RobustRA()`**: It executes the Robust Rank Aggregation (RRA) method of Kolde et al., 2012.

- **void `DIBRA()`**: It executes the distance-based iterative rank aggregation method of Akritidis et al., 2022.

- **void `PrefRel()`**: It executes the preference relations method of Desarkar et al., 2016.

- **void `Agglomerative()`**: It executes the agglomerative rank aggregation method of Chatterjee et al., 2018.

- **void `MC()`**: It executes the Markov Chain-based rank aggregation method of Dwork et al., 2001.

# 1.5 Execution driver

The Execution Driver is a very simple function defined in `flagr.cpp`. Its role is to offer a unified manner of executing the exposed C functions. It takes as an argument a simple C structure that stores the user-defined input parameters and algorithm hyper-parameters, and orchestrates the execution flow.

More specifically, the Execution driver initially copies the user-defined input parameters from the aforementioned C structure to an `InputParams` C++ object. In the sequel, it creates an `InputData` object which immediately starts reading the provided input file/s. Next, the `aggregate()` method of `InputData` is called to perform rank aggregation. Notice that the `InputParams` object is propagated to all FLAGR components and carries all the required parameters with it. Therefore, the rank aggregation method to be applied is automatically executed without any further checks, since `InputParams` notifies the `Aggregator` about the aggregation method that was selected by the user.

In case the user provided a valid input file with relevance judgments, the Execution Driver proceeds to the evaluation of the generated aggregate list. This is achieved by calling the `evaluate()` method of `InputData`, which in turn triggers the `Evaluator`.

The following block diagram depicts the role of the execution driver as the connector between the dynamic library references and the C++ FLAGR core.

# 2. Rank Aggregation methods

This section describes the rank aggregation algorithms implemented in FLAGR. The following articles also contain brief descriptions of the respective exposed C functions that make these algorithms accessible from other programs (in case the FLAGR shared library is used).

## 2.1 Linear Combination methods – The `Linear()` function

In linear combination methods, the score of each element is computed by summing up the partial scores of that element with respect to its rankings in each input preference list. The `Linear()` function triggers the execution of two such combination methods: CombSUM, and CombMNZ. Both of them are implemented in accordance to the paper of Renda et al., 2003.

Each method is accompanied by an element rank/score normalization technique, as it is described in the aforementioned paper. These techniques are: Rank normalization, Borda normalization, Score normalization, and Z-Score normalization. In FLAGR, there is a fifth normalization technique, called Simple Borda. In contrast to the traditional Borda normalization, Simple Borda assigns zero partial scores in case an element has not been ranked by an input preference list.

In addition, notice that CombSUM with Borda normalization is equivalent to the well-known BordaCount rank aggregation method.

**Function Definitions**

```
void Linear(const char inf[], const char relf[], const int evpts, const int
ram, const char ranstr[], const char out[])
```

and

```
__declspec(dllexport) void __cdecl Linear(const char inf[], const char relf[],
const int evpts, const int ram, const char ranstr[], const char out[])
```

**Implementation Files**

- `Linear()` function: `cflagr.cpp` and `dllflagr.cpp`.
- CombSUM: `src/ram/CombSUM.cpp`.
- CombMNZ: `src/ram/CombMNZ.cpp`.

**Input arguments**

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.
- `const char relf[]`: The path of the input file that stores the element relevance judgments.
- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- `const int ram`: The selected rank aggregation method (see the `aggregation_method` parameter in this document for a list of possible values).
- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.

- `const char out[]`: The file system location where the output file with the aggregate list will be stored.

### Description

The input parameters are parsed and stored in a special C structure called `UserParams` that is defined in `src/InputParams.h`. Then, `UserParams` is passed to the execution driver and the rank aggregation process starts.

# 2.2 Condorcet Winners – The `Condorcet()` function

The `Condorcet()` function executes the Condorcet Winners method. The score of an element $r_i$ is determined by the number of its "victories" against all the other involved elements. A victory for $r_i$ is achieved if the majority of the voters rank $r_i$ higher than another element $r_j$. Finally, the elements are sorted in decreasing victory order.

### Function Definitions

```
void Condorcet(const char inf[], const char relf[], const int evpts, const char
ranstr[], const char out[])
```

and

```
__declspec(dllexport) void __cdecl Condorcet(const char inf[], const char
relf[], const int evpts, const char ranstr[], const char out[])
```

### Implementation Files

- `Condorcet()` function: `cflagr.cpp` and `dllflagr.cpp`.
- Condorcet Winners method: `src/ram/CondorcetWinners.cpp`.

### Input arguments

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.
- `const char relf[]`: The path of the input file that stores the element relevance judgments.
- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.
- `const char out[]`: The file system location where the output file with the aggregate list will be stored.

### Description

The input parameters are parsed and stored in a special C structure called `UserParams` that is defined in `src/InputParams.h`. Then, `UserParams` is passed to the execution driver and the rank aggregation process starts.

## 2.3 Copeland Winners – The `Copeland()` function

The `Copeland()` function executes the Copeland Winners method. The score of an element $r_i$ is determined by the number of its "victories" against all the other involved elements. A victory for $r_i$ is achieved if the majority of the voters rank $r_i$ higher than another element $r_j$. In contrast to the Condorcet method, Copeland Winners assign "half" a victory (i.e., a score of 0.5) to each element of a pair $(r_i, r_j)$ in the case of tie. A tie happens when exactly half of the voters rank $r_i$ higher than $r_j$ and the other half voters rank $r_j$ higher than $r_i$.

Finally, the elements are sorted in decreasing victory order.

### Function Definitions

```
void Copeland(const char inf[], const char relf[], const int evpts, const char ranstr[], const char out[])
```

and

```
__declspec(dllexport) void __cdecl Copeland(const char inf[], const char relf[], const int evpts, const char ranstr[], const char out[])
```

### Implementation Files

- `Copeland()` function: `cflagr.cpp` and `dllflagr.cpp`.
- Copeland Winners method: `src/ram/CopelandWinners.cpp`.

### Input arguments

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.
- `const char relf[]`: The path of the input file that stores the element relevance judgments.
- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.
- `const char out[]`: The file system location where the output file with the aggregate list will be stored.

### Description

The input parameters are parsed and stored in a special C structure called `UserParams` that is defined in `src/InputParams.h`. Then, `UserParams` is passed to the execution driver and the rank aggregation process starts.


## 2.4 Outranking Approach – `OutrankingApproach()`

The Outranking Approach of Farah & Vanderpooten is a majoritarian method that identifies the "winning" elements by performing pairwise comparisons of their individual rankings. The method is implemented in accordance to the following paper:

- Farah, M., Vanderpooten, D., "An outranking approach for rank aggregation in information retrieval", In Proceedings of the 30th ACM Conference on Research and Development in Information Retrieval, pp. 591-598, 2007.

The algorithm is based on four threshold values that introduce different perspectives of the majority criterion. These values are the concordance, discordance, preference, and veto thresholds. The user may pass all of them to FLAGR as hyper-parameters, through the input arguments of the `OutrankingApproach()` function (see below).

## Function Definitions

```
void OutrankingApproach(const char inf[], const char relf[], const int evpts,
const char ranstr[], const char out[], const float pref_t, const float veto_t,
const float conc_t, const float disc_t)
```

and

```
__declspec(dllexport) void __cdecl OutrankingApproach(const char inf[], const
char relf[], const int evpts, const char ranstr[], const char out[], const float
pref_t, const float veto_t, const float conc_t, const float disc_t)
```

## Implementation Files

- The `OutrankingApproach()` C function: `cflagr.cpp` and `dllflagr.cpp`.
- Algorithm Implementation: `src/ram/OutrankingApproach.cpp`.

## Input arguments

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.
- `const char relf[]`: The path of the input file that stores the element relevance judgments.
- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.
- `const char out[]`: The file system location where the output file with the aggregate list will be stored.
- `const float pref_t`: Algorithm hyper-parameter - The value of the preference threshold.
- `const float veto_t`: Algorithm hyper-parameter - The value of the veto threshold.
- `const float conc_t`: Algorithm hyper-parameter - The value of the concordance threshold.
- `const float disc_t`: Algorithm hyper-parameter - The value of the discordance threshold.

## Description

The input parameters are parsed and stored in a special C structure called `UserParams` that is defined in `src/InputParams.h`. Then, `UserParams` is passed to the execution driver and the rank aggregation process starts.

# 2.5 Markov Chains methods – The `MC()` function

The Markov Chains methods constitute a well-established family of rank aggregation methods. Originally proposed by Dwork et al., (2001), they consider an aggregate list as a system that moves from one state to another with respect to a particular criterion. Dwork et al. (2001) introduced four such methods in the following article:

- C. Dwork, R. Kumar, M. Naor, D. Sivakumar, "Rank Aggregation Methods for the Web", In Proceedings of the 10th International Conference on World Wide Web, pp. 613-622, 2001.

In addition, DeConde et al. (2006) introduced MCT, a variant that constructs the transition matrix by considering the proportion of lists which prefer an element over another.

- R.P. DeConde, S. Hawley, S. Falcon, N. Clegg, B. Knudsen, R. Etzioni, "Combining results of microarray experiments: a rank aggregation approach", Statistical Applications in Genetics and Molecular Biology, vol. 5, no. 1, 2006.

The execution of all five methods takes place by passing the appropriate arguments to the MC exposed C function of FLAGR.

## Function Definitions

```
void MC(const char inf[], const char relf[], const int evpts, const int ram,
const char ranstr[], const char out[], const float ergodic_number, const float
delta, const int iter)
```

and

```
__declspec(dllexport) void __cdecl MC(const char inf[], const char relf[], const
int evpts, const int ram, const char ranstr[], const char out[], const float
ergodic_number, const float delta, const int iter)
```

## Implementation Files

- `MC()` function: `cflagr.cpp` and `dllflagr.cpp`.
- Markov Chains implementation: `src/ram/MC.cpp`.

## Input arguments

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.
- `const char relf[]`: The path of the input file that stores the element relevance judgments.
- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- `const int ram`: The selected (Markov Chains-based) rank aggregation method (801, 802, 803, 804, or 805 - see the `aggregation_method` parameter in this document).
- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.
- `const char out[]`: The file system location where the output file with the aggregate list will be stored.
- `const float ergodic_number`: The ergodic number, used during the computation of the ergodic transition matrix from the normalized transition matrix.

- **const float iter**: Controls the maximum number of iterations before convergence.

### Description

The input parameters are parsed and stored in a special C structure called UserParams that is defined in src/InputParams.h. Then, UserParams is passed to the execution driver and the rank aggregation process starts.

# 2.6 Kemeny Optimal Aggregation – The Kemeny() function

This function executes Kemeny optimal aggregation. This algorithm identifies the optimal aggregate list as the list that minimizes its distance from all the input preference lists.

Kemeny optimal aggregation is an NP-hard problem, with very high computational complexity. It requires the computation of all permutations of the input items and the calculation of the distances of each permutation from all input lists. *The brute force solution becomes infeasible when the number of elements gets greater than 15-20, or the number of input lists is greater than 4*, so caution is advised.

### Function Definitions

```
void Kemeny(const char inf[], const char relf[], const int evpts, const char
ranstr[], const char out[])
```

and

```
__declspec(dllexport) void __cdecl Kemeny(const char inf[], const char relf[],
const int evpts, const char ranstr[], const char out[])
```

### Implementation Files

- Kemeny() function: cflagr.cpp and dllflagr.cpp.
- Kemeny optimal aggregation method: src/ram/KemenyOptimal.cpp.

### Input arguments

- **const char inf[]**: The path of the input file that stores the preference lists to be aggregated.
- **const char relf[]**: The path of the input file that stores the element relevance judgments.
- **const int evpts**: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- **const char ranstr[]**: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.
- **const char out[]**: The file system location where the output file with the aggregate list will be stored.

### Description

The input parameters are parsed and stored in a special C structure called UserParams that is defined in src/InputParams.h. Then, UserParams is passed to the execution driver and the rank aggregation process starts.

# 2.7 Robust Rank Aggregation – The `RobustRA()` function

This function executes the Robust Rank Aggregation (RRA) method of Kolde et al., 2012. The method is implemented in accordance to the following paper:

- R. Kolde, S. Laur, P. Adler, J. Vilo, "Robust rank aggregation for gene list integration and meta-analysis", Bioinformatics, vol. 28, no. 4, pp. 573-580, 2012.

RRA is mostly used in bio-informatics applications to aggregate gene lists. It is based on a probabilistic model (beta distribution) that makes the algorithm parameter free and robust to outliers, noise and errors. The FLAGR C++ implementation of RRA produces the same results as the R implementation of Kolde (see the RobustRankAggreg R package).

The computation of the incomplete beta function is performed with the John Burkardt's implementation of ASA063 algorithm (K.L. Majumder and G. Bhattacharjee):

- K.L. Majumder, G.P. Bhattacharjee, "Algorithm AS 63: The incomplete Beta Integral", Applied Statistics, vol. 22, no. 3, pp. 409-411, 1973.

Furthermore, the computation of the inverse of the incomplete beta function is performed with the John Burkardt's implementation of ASA109 algorithm (GW Cran, KJ Martin and GE Thomas):

- G.W. Cran, M.J. Martin, G.E. Thomas, "Remark AS R19 and Algorithm AS 109: A Remark on Algorithms AS 63: The Incomplete Beta Integral and AS 64: Inverse of the Incomplete Beta Integeral", Applied Statistics, Volume 26, Number 1, 1977, pages 111-114.

## Function Definitions

```
void RobustRA(const char inf[], const char relf[], const int evpts, const char
ranstr[], const char out[], const bool exact)
```

and

```
__declspec(dllexport) void __cdecl RobustRA(const char inf[], const char relf[],
const int evpts, const char ranstr[], const char out[], const bool exact)
```

## Implementation Files

- `RobustRA()` function: `cflagr.cpp` and `dllflagr.cpp`.
- RRA implementation: `src/ram/RobustRA.cpp`.
- Incomplete beta function implementation (ASA 063 algorithm), inverse of the incomplete Beta function implementation (ASA 109 algorithm): `src/ram/tools/BetaDistribution.cpp`.

## Input arguments

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.
- `const char relf[]`: The path of the input file that stores the element relevance judgments.
- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.

- `const char out[]`: The file system location where the output file with the aggregate list will be stored.
- `const bool exact`: If true the Stuart algorithm for p-value correction is applied.

**Description**

The input parameters are parsed and stored in a special C structure called `UserParams` that is defined in `src/InputParams.h`. Then, `UserParams` is passed to the execution driver and the rank aggregation process starts.

# 2.8 Iterative Distance-Based weighted aggregation – The `DIBRA()` function

This function executes the iterative, distance-based method (abbreviated DIBRA) of Akritidis el. al 2022. The method is implemented in accordance to the following paper:

- L. Akritidis, A. Fevgas, P. Bozanis, Y. Manolopoulos, "An Unsupervised Distance-Based Model for Weighted Rank Aggregation with List Pruning", Expert Systems with Applications, vol. 202, pp. 117435, 2022.

DIBRA belongs to the weighted rank aggregation methods. It employs exploratory analysis to automatically identify the expert voters in an unsupervised fashion. Then, it assigns higher weights to the voters who were identified as experts, thus boosting the scores of their submitted elements.

In particular, DIBRA employs a standard non-weighted method to generate an initial aggregate ranking (see `aggregation_method` in this document for a list of the supported methods). Then, it repeatedly assigns higher weights to the input lists that have smaller distances from the aggregate lists. The process terminates when the voter weights converge and a stable aggregate list is obtained.

The algorithm also includes an optional list pruning mechanism that arranges the input list lengths according to the respective voter weights.

**Function Definitions**

```
void DIBRA(const char inf[], const char relf[], const int evpts, const int agg,
const char ranstr[], const char out[], const int wnorm, const int dist, const
bool prune, const float gamma, const float d1, const float d2, const float tol,
const int iter, const float pref_t, const float veto_t, const float conc_t,
const float disc_t)
```

and

```
__declspec(dllexport) void __cdecl DIBRA(const char inf[], const char relf[],
const int evpts, const int agg, const char ranstr[], const char out[], const
int wnorm, const int dist, const bool prune, const float gamma, const float d1,
const float d2, const float tol, const int iter, const float pref_t, const float
veto_t, const float conc_t, const float disc_t)
```

**Implementation Files**

- `DIBRA()` function: `cflagr.cpp` and `dllflagr.cpp`.
- DIBRA method implementation: `src/ram/DIBRA.cpp`.

**Input arguments**

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.

- `const char relf[]`: The path of the input file that stores the element relevance judgments.

- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.

- `const int agg`: The selected non-weighted base rank aggregation method (see the aggregation_method parameter in this document for a list of possible values).

- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.

- `const char out[]`: The file system location where the output file with the aggregate list will be stored.

- `const int wnorm`: The voter weights normalization method (see the weights_normalization parameter in this document for a list of possible values).

- `const int dist`: The correlation method that is used to measure the distance between an input list and the temporary aggregate list (see the `correlation_method` parameter in this document for a list of possible values).

- `const bool prune`: Triggers a weight-dependent list pruning mechanism.

- `const float gamma`: The $\gamma$ hyper-parameter.

- `const float d1`: The $\delta_1$ hyper-parameter (applicable when prune=true).

- `const float d2`: The $\delta_2$ hyper-parameter (applicable when prune=true).

- `const float tol`: Controls the convergence precision. This tolerance threshold represents the minimum precision of the difference between the voter weight in an iteration and the voter weight of the previous iteration.

- `const int iter`: Controls the maximum number of iterations.

- `const float pref_t`: The preference threshold (applicable when the Outranking Approach is selected as the non-weighted base method; namely, `agg=5300`).

- `const float veto_t`: The veto threshold (applicable when the Outranking Approach is selected as the non-weighted base method; namely, `agg=5300`).

- `const float conc_t`: The concordance threshold (applicable when the Outranking Approach is selected as the non-weighted base method; namely, `agg=5300`).

- `const float disc_t`: The discordance threshold (applicable when the Outranking Approach is selected as the non-weighted base method; namely, `agg=5300`).

**Description**

The input parameters are parsed and stored in a special C structure called `UserParams` that is defined in `src/InputParams.h`. Then, `UserParams` is passed to the execution driver and the rank aggregation process starts.

# 2.9 Preference Relations weighted method

This function executes the Preference Relations weighted rank aggregation method of Desarkar et al., 2016. The method is implemented in accordance to the following paper:

- M.S. Desarkar, S. Sarkar, P. Mitra, "Preference relations based unsupervised rank aggregation for metasearch", Expert Systems with Applications, vol. 49, pp. 86-98, 2016.

The Preference Relations algorithm belongs to the weighted rank aggregation methods. It employs exploratory analysis to automatically identify the expert voters in an unsupervised fashion. Then, it assigns higher weights to the voters who were identified as experts, thus boosting the scores of their submitted elements.

The method constructs a preference relations graph which contains the individual elements as vertices and their weights as edges.

## Function Definitions

```
void PrefRel(const char inf[], const char relf[], const int evpts, const char ranstr[], const char out[], const float alpha, const float beta)
```

and

```
__declspec(dllexport) void __cdecl PrefRel(const char inf[], const char relf[], const int evpts, const char ranstr[], const char out[], const float alpha, const float beta)
```

## Implementation Files

- PrefRel() function: cflagr.cpp and dllflagr.cpp.
- Preference Relations method implementation: src/ram/PrefRel.cpp.
- Helper class MergedItemPair implementation: src/ram/tools/MergedItemPair.cpp.

## Input arguments

- const char inf[]: The path of the input file that stores the preference lists to be aggregated.
- const char relf[]: The path of the input file that stores the element relevance judgments.
- const int evpts: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- const char ranstr[]: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.
- const char out[]: The file system location where the output file with the aggregate list will be stored.
- const float alpha: The $\alpha$ hyper-parameter.
- const float beta: The $\beta$ hyper-parameter.

## Description

The input parameters are parsed and stored in a special C structure called UserParams that is defined in src/InputParams.h. Then, UserParams is passed to the execution driver and the rank aggregation process starts.

# 2.10 Agglomerative weighted aggregation

This function executes the Agglomerative weighted rank aggregation method of Chatterjee et al., 2018. The method is implemented in accordance to the following paper:

- S. Chatterjee, A. Mukhopadhyay, M. Bhattacharyya, "A weighted rank aggregation approach towards crowd opinion analysis", Knowledge-Based Systems, vol. 149, pp. 47-60, 2018.

The Agglomerative Aggregation algorithm belongs to the weighted rank aggregation methods. It employs exploratory analysis to automatically identify the expert voters in an unsupervised fashion. Then, it assigns higher weights to the voters who were identified as experts, thus boosting the scores of their submitted elements.

This method works very similarly to the well-established agglomerative clustering algorithm. Specifically, it repeatedly merges the two most similar input lists into a temporary aggregate list. During list merging, it modifies the weights of the respective voters, thus affecting the future merges.

## Function Definitions

```
void Agglomerative(const char inf[], const char relf[], const int evpts, const
char ranstr[], const char out[], const float c1, const float c2)
```

and

```
__declspec(dllexport) void __cdecl Agglomerative(const char inf[], const char
relf[], const int evpts, const char ranstr[], const char out[], const float c1,
const float c2)
```

## Implementation Files

- `Agglomerative()` function: `cflagr.cpp` and `dllflagr.cpp`.
- Agglomerative Aggregation method implementation: `src/ram/Agglomerative.cpp`.

## Input arguments

- `const char inf[]`: The path of the input file that stores the preference lists to be aggregated.
- `const char relf[]`: The path of the input file that stores the element relevance judgments.
- `const int evpts`: the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed.
- `const char ranstr[]`: A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library.
- `const char out[]`: The file system location where the output file with the aggregate list will be stored.
- `const float c1`: The $c_1$ hyper-parameter.
- `const float c2`: The $c_2$ hyper-parameter.

## Description

The input parameters are parsed and stored in a special C structure called `UserParams` that is defined in `src/InputParams.h`. Then, `UserParams` is passed to the execution driver and the rank aggregation process starts.

# 3. FLAGR API Documentation

In this section the FLAGR library core is presented. The following Figure depicts a block diagram of the core's architecture.



In the following subsections, the format of the input and output files is described. Then, the participating classes of FLAGR are presented. The final Subsection contains a detailed guide on how custom algorithm implementations can be integrated to FLAGR.

## 3.1 Input and Output files

### Input 1: The file of the input preference lists

FLAGR requires that the input preference lists to be aggregated are stored in a single CSV file, regardless of the number of the involved topics (queries) or voters (rankers). The columns of this CSV file must be organized in the following manner:

```
Query/Topic String, Voter Name, Item Code, Item Score, Algorithm/Dataset
```

where:

- `Query/Topic`: the query string or the topic for which the preference list is submitted.

- `Voter`: the name of the voter, or the ranker who submits the preference list for the specified Query/Topic.

- `Item Code`: a unique name that identifies a particular element of the preference list. A voter cannot submit the same element for the same query/topic two or more times. This means that each element appears exactly once in each preference list. However, the same element may appear in lists submitted by other voters.

- `Item Score`: the preference score assigned to an item by a specific voter. It reflects the importance (or the relevance, or the weight) of the element. In many cases (e.g., search engine rankings), the preference scores are unknown. In such cases, the scores can be replaced by the (reverse) ranking of an item in such a manner that the top rankings receive higher scores than the ones that have been assigned lower rankings.

- `Algorithm/Dataset`: A user-defined string that usually represents the origin of a particular preference list. It may receive any non-blank value.

You may find an example of an input list CSV file here. This example file contains the preference lists that were submitted by 50 voters for 20 queries. Each input list contains 30 elements. Therefore, the number of rows in this file is equal to 50×20×30=30000.

## Output 1: The file of the aggregate list/s

In this file FLAGR stores the result (output) of the selected rank aggregation method. Namely, the final lists that derive after the aggregation of the input preference lists. The library creates one aggregate list per input query/topic. So, if there are $Q$ input queries, FLAGR generates $Q$ aggregate lists and stores them in a CSV file. Each row in the file represents an element of the aggregate list stored in decreasing score order. The columns are organized as follows:

```
Query/Topic String, Voter Name, Item Code, Item Score
```

## Input 2: The file of relevant elements (or, the Rels file)

Optionally, the user may provide a second CSV file (we call it Rels file) that contains relevance judgments for the preference list elements of the primary input file for each query. The Rels file is employed by the FLAGR's evaluation module to evaluate each created aggregate list. Its columns must be formatted as follows:

```
Query/Topic String, 0, Item Code, Relevance Score
```

where:

- `Query/Topic`: the query string or the topic for which the list is submitted.

- `0`: unused. This value *must be always* 0.

- `Item Code`: a unique name that identifies a particular element. There cannot be two relevance judgments for the same element for the same query.

- `Relevance Score`: An integer value that represents the relevance of the item with respect to the mentioned Query/Topic. Typically, zero values represent irrelevant and incorrect elements, negative values represent spam elements; and positive values represent relevant, correct and informative elements.

You may find an example of an input Rels file here. This example file contains the relevance judgments for the elements of all preference lists for all queries of the previous input list file. *Notice that in case FLAGR does not find a relevance judgment for an element, then it automatically considers it as irrelevant (that is, it sets its Relevance Score equal to 0).*

## Output 2: The evaluation file

As soon as a valid Rels file is provided, the evaluation process takes place automatically. In this case FLAGR evaluates each aggregate list individually and outputs a second CSV file, where it writes the results of the evaluation.

If there are $Q$ input queries, then $Q$ aggregate lists are generated and the evaluation file contains $Q + 1$ rows. The first $Q$ rows store the evaluation metrics for each aggregate list, whereas the last row contains the average values. On the other hand, the columns of the evaluation file depend on the eval_pts parameter that is set by the user. More specifically, the columns are:

```
Average Precision, P@1, P@2, ..., P@eval_pts, N@1, N@2, ..., N@eval_pts
```

where

- Average Precision is the Average Precision for this aggregate list.

- P@X is the running Precision at the X-th element of the aggregate list.

- N@X is the running Normalized Discounted cumulative gain (NDCG) at the X-th element of the aggregate list.

# 3.2 Aggregator class

The Aggregator class triggers the execution of a rank aggregation algorithm on a collection (array) of InputLists. Typically, the output of the aggregation process is a single MergedList object.

**Implementation files**

The Aggregator class is defined in the src/Aggregator.h header file; its member functions are implemented in src/Aggregator.cpp.

**Technical Details**

The input_lists member variable stores the input lists to be aggregated. Technically, it is an array (i.e. double pointer) of InputList objects. The allocated size of this array is equal to num_alloc_lists, whereas the real number of non-null input lists is num_lists. Naturally, it must always hold that num_alloc_lists ≥ num_lists.

The rank aggregation process takes places inside a Query object. For this reason, a Query object contains a pointer to a single Aggregator object (see the respective block diagram).

The aggregate() member function is responsible for executing the rank aggregation procedure. It accepts an InputParams object that stores the selected rank aggregation method, the hyper-parameter values and the execution parameters and returns a single MergedList.

# 3.3 Evaluator class

The Evaluator class evaluates the quality of an aggregate list with respect to a given set of judgments that determine the relevance of (some, or all) list elements. The set of relevance judgments is stored in a Rels object. The results of the evaluation (i.e. evaluation measures) are written in a CSV file.

**Implementation files**

The Evaluator class is defined in the src/Evaluator.h header file; its member functions are implemented in src/Evaluator.cpp.

**Technical Details**

The evaluation process takes place inside a Query object, provided that a file of relevance judgments is provided to FLAGR (see more in the article about Input and Output files). For this reason, a Query object contains a pointer to a single `Evaluator` object.

The evaluation process is performed by the `evaluate()` member function. The procedure is based on the Rels object that contains the aforementioned relevance judgments. Finally, the four following arrays are created:

- `precision`: in its $i$-th position, it stores the value of the Precision measure at the $i$-th element of the aggregate list.

- `recall`: in its $i$-th position, it stores the value of the Recall measure at the $i$-th element of the aggregate list.

- `dcg`: in its $i$-th position, it stores the value of the Discounted Cumulative Gain (DCG) measure at the $i$-th element of the aggregate list.

- `ndcg`: in its $i$-th position, it stores the value of the normalized Discounted Cumulative Gain (nDCG) measure at the $i$-th element of the aggregate list.

# 3.4 `InputItem` class

An `InputItem` is a single element that is read from the input file. It is a part of an `InputList` and possesses three properties:

- A unique string identifier that is used by FLAGR to identify the common elements among all `InputLists`,

- Its (integer) ranking in the input preference list, and

- An optional score value that is assigned by the associated voter and justifies its ranking in the input list. The score data type can be either `float`, or `double`; this is determined by the `score_t` data type definition in `flagr.cpp`.

**Implementation files**

The `InputItem` class is defined in the `src/InputItem.h` header file; its member functions are implemented in `src/InputItem.cpp`.

# 3.5 `InputList` class

This class is used to store and represent an input preference list submitted by a voter. It consists of an array of `InputItem` objects and it is connected to the voter who submitted it via a pointer that points to a Voter object.

The input lists are read from the input file by using an `InputData` object. The entire collection of them is handled by an Aggregator.

**Implementation files**

The `InputList` class is defined in the `src/InputList.h` header file; its member functions are implemented in `src/InputList.cpp`.

**Technical Details**

The actual number of elements in an `InputList` is `num_items`; the allocated memory is `num_alloc_items`. Naturally it derives that `num_alloc_items` must always be greater than, or equal to `num_items`.

# 3.6 `InputParams` class

The `InputParams` class stores options and execution parameters that have been passed to FLAGR by the user. These parameters concern input and output file locations, rank aggregation methods, algorithm hyper-parameters, etc. See the table below for a complete list of the supported parameters and their respective valid values.

This object is passed as an argument to multiple functions of FLAGR including the implementations of the rank aggregation methods. This is how these implementations get access to the user-defined hyper-parameters.

**Implementation files**

The `InputParams` class is defined in the `src/InputParams.h` header file; its member functions are implemented in `src/InputParams.cpp`.

**Details**

The supported parameters include:

| Parameter | Data Type | Description |
|---|---|---|
| `input_file` | String (ASCII) | The path of the input CSV file that contains the preference lists to be aggregated. |
| `rels_file` | String (ASCII) | The path of the optional CSV file that contains the relevance judgments of the input list elements. If set, it automatically triggers the evaluation process of the generated aggregate list. Otherwise, no evaluation takes place. |
| `output_file` | String (ASCII) | The file system location where the output file with the aggregate list will be stored. If left empty, the default OS temp directory is used. |
| `eval_file` | String (ASCII) | The file system location where the output file with the evaluation of the aggregate list will be stored. If left empty, the default OS temp directory is used. |
| `random_string` | String (ASCII) | A string that is embedded in the names of the output files. Used when FLAGR is compiled as a shared library. |

| | | |
|---|---|---|
| aggregation_method | Integer | Determines the algorithm that will be used to perform rank aggregation. The valid values are:<br><br>100: for CombSUM with Borda normalization<br>101: for CombSUM with Rank normalization<br>102: for CombSUM with Score normalization<br>103: for CombSUM with Z-Score normalization<br>104: for CombSUM with Simple Borda normalization<br>110: for CombMNZ with Borda normalization<br>111: for CombMNZ with Rank normalization<br>112: for CombMNZ with Score normalization<br>113: for CombMNZ with Z-Score normalization<br>114: for CombMNZ with Simple Borda normalization<br>200: for the Condorcet Winners method<br>201: for the Copeland Winners method<br>300: for the Outranking Approach<br>400: for Kemeny optimal aggregation<br>401: for Robust Rank Aggregation (RRA)<br>5100: for DIBRA with CombSUM and Borda Normalization<br>5101: for DIBRA with CombSUM and Rank Normalization<br>5102: for DIBRA with CombSUM and Score Normalization<br>5103: for DIBRA with CombSUM and Z-Score Normalization<br>5104: for DIBRA CombSUM and Simple Borda Normalization<br>5110: for DIBRA with CombMNZ and Borda Normalization<br>5111: for DIBRA with CombMNZ and Rank Normalization<br>5112: for DIBRA with CombMNZ and Score Normalization<br>5113: for DIBRA with CombMNZ and Z-Score Normalization<br>5114: for DIBRA CombMNZ and Simple Borda Normalization<br>5200: for DIBRA with the Condorcet Winners method<br>5201: for DIBRA with the Copeland Winners method<br>5300: for DIBRA with the Outranking Approach<br>600: for the Preference Relations Method<br>700: for the Agglomerative Aggregation Method<br>801: for Markov Chains 1 (MC1)<br>802: for Markov Chains 2 (MC2)<br>803: for Markov Chains 3 (MC3)<br>804: for Markov Chains 4 (MC4)<br>805: for MCT |
| correlation_method | Integer | The correlation method that is used to measure the distance between an input list and the temporary aggregate list. The valid values are:<br><br>1: for the Spearman's $\rho$ correlation coefficient.<br>2: for the scaled variant of Spearman's Footrule distance.<br>3: for Cosine similarity of the lists' vector representations.<br>5: for the Kendall's $\tau$ correlation coefficient. |

| | | |
|---|---|---|
| `weights_normalization` | Integer | The voter weights normalization method. Used when the DIBRA algorithm is selected. The valid values are:<br>1: for no voter weight normalization.<br>2: for normalizing the voter weights with min-max scaling.<br>3: for z-normalizing the voter weights. |
| `max_iterations` | Integer | This parameter controls the maximum number of iterations. FLAGR will stop the execution of DIBRA if the requested number of iterations have been performed, even if the voter weights have not fully converged. |
| `max_list_items` | Integer | Limits the length of the input preference lists. |
| `eval_points` | Integer | Determines the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute $P@1, P@2, \ldots, P@10$ and $N@1, N@2, \ldots, N@10$. |
| `list_pruning` | Boolean | Triggers a weight-dependent list pruning mechanism. Used in combination with the DIBRA weighted method only. |
| `convergence_precision` | Float or Double | Controls the convergence precision. This tolerance threshold represents the minimum precision of the difference between the voter weight in an iteration and the voter weight of the previous iteration. Used in combination with the DIBRA weighted method only. |
| `alpha` | Float or Double | The $\alpha$ hyper-parameter of the Preference Relations method. |
| `beta` | Float or Double | The $\beta$ hyper-parameter of the Preference Relations method. |
| `gamma` | Float or Double | The $\gamma$ hyper-parameter of DIBRA. |
| `c1` | Float or Double | The $c_1$ hyper-parameter of the Agglomerative Aggregation method. |
| `c2` | Float or Double | The $c_2$ hyper-parameter of the Agglomerative Aggregation method. |
| `pref_thr` | Float or Double | The preference threshold of the Outranking Approach. It takes values in the range [0,1]. |
| `veto_thr` | Float or Double | The veto threshold of the Outranking Approach. It takes values in the range [0,1]. |
| `conc_thr` | Float or Double | The concordance threshold of the Outranking Approach. It takes values in the range [0,1]. |
| `disc_thr` | Float or Double | The discordance threshold of the Outranking Approach. It takes values in the range [0,1]. |

## 3.7 `InputData` class

This class is responsible for reading and parsing the input data file/s. For the moment, FLAGR accepts only CSV-formatted input files.

**Implementation files**

The `InputData` class is defined in the `src/input/InputData.h` header file; its member functions are implemented in `src/input/InputData.cpp` and `src/input/InputDataCSV.cpp`.

**Input list files**

Detailed information about how the input files must be formatted can be found here.

Also notice that FLAGR is designed to accept data directly from RASDaGen, a synthetic dataset generator for rank aggregation problems.

**Technical Details**

The role of `InputData` is broader and it is not limited to just reading the input files. More specifically, during the input file parsing process, one or more `Query` objects are constructed. In the sequel, inside each `Query` the corresponding `Aggregator` and `Evaluator` objects are created to initialize and evaluate the aggregation process.

There is also a pointer called `params` that connects `InputData` with `InputParams`. In this manner, `InputData` is able to subsequently pass user-defined algorithm hyper-parameters and several other execution parameters to the rest of the application components.


## 3.8 `MergedItem` class

A `MergedItem` is an element of an aggregate `MergedList`. The class derives from `InputItem` and inherits all its member variables and functions.

**Implementation files**

The MergedItem class is defined in the `src/MergedItem.h` header file; its member functions are implemented in `src/MergedItem.cpp`.

**Technical Details**

`MergedItem` maintains an array of `Ranking` objects that store the individual rankings (and scores) of the element in each preference `InputList`. Notice that the number of elements of this array (`num_alloc_rankings`) is always equal to the number of the input preference lists. In case the associated `InputItem` has not been ranked by a preference list (i.e. it is not included in the list), then its corresponding ranking in the `Rankings` array is set equal to `NOT_RANKED_ITEM_RANK` (defined in `flagr.cpp`). The actual number of the input preference lists that include a particular `MergedItem` is stored in the `num_rankings` member variable. This approach consumes more memory, but enables $O(1)$ constant search times of the ranking of a particular item on a particular `InputList`.

The class also includes a self-class next pointer that points to another `MergedItem` object. This allows the storage of a collection of `MergedItems` in dynamic data structures, e.g. linked lists. Such structures are employed in `MergedList`, where `MergedItems` are stored in a hash table with linked lists as chains.

# 3.9 `MergedList` class

The result of the rank aggregation process is an aggregate list that is stored in a `MergedList` object. `MergedList` contains a collection of `MergedItems`, typically sorted in decreasing score order. The score is assigned by a rank aggregation method.

Regarding the evaluation, the generated aggregate list is fed to the `evaluate()` function of an `Evaluator`. The results are written in a CSV file according to this document.

**Implementation files**

The `MergedList` class is defined in the `src/MergedList.h` header file; its member functions are implemented in `src/MergedList.cpp`. The implementations of the supported rank aggregation methods are stored in the `/ram` directory. More specifically:

- `src/ram/Agglomerative.cpp` implements the Agglomerative Aggregation method of Chatterjee et al., 2018.

- `src/ram/CombMNZ.cpp` implements the CombMNZ linear combination methods as they are described in the paper of Renda et al., 2003.

- `src/ram/CombSUM.cpp` implements the CombSUM linear combination methods (including Borda Count) as they are described in the paper of Renda et al., 2003.

- `src/ram/CondorcetWinners.cpp` implements one of the oldest approaches to rank aggregation, the Condorcet criterion method.

- `src/ram/CopelandWinners.cpp` implements the method of Copeland Winners which is a variant of the Condorcet Winners method.

- `src/ram/CustomMethods.cpp` contains two sample functions that facilitate the integration of custom rank aggregation methods.

- `src/ram/DIBRA.cpp` implements the Iterative Distance-Based Weighted method of Akritidis et al., 2022.

- `src/ram/KemenyOptimal.cpp` contains the brute force implementation of Kemeny optimal aggregation.

- `src/ram/MC.cpp` implements the four Markov Chains method of Dwork et al., 2001 and the MCT method of DeConde et al., 2006.

- `src/ram/OutrankingApproach.cpp` implements the Outranking Approach of Farah and Vanderpooten, 2007.

- `src/ram/PrefRel.cpp` implements the Preference Relations Weighted method of Desarkar et al., 2016.

- `src/ram/RobustRA.cpp` implements the Robust Rank Aggregation method (RRA) of Kolde et al., 2012.

For more details, please visit the Publications section, or follow the links in the Introduction section. A guide on how custom rank aggregation implementations can be integrated to FLAGR is given here.

**Technical Details**

The elements of `MergedList` are organized in two ways. More specifically:

- A hash table (`hash_table` member variable) with linked lists as chains (for collision resolution) is employed to support fast fusion of the individual `InputLists`. The contents of this hash table are `MergedItem` objects; the search keys are the unique identifiers of the associated `InputItems` (`MergedItem` inherits the members of `InputItem`).

- A typical array of `MergedItem` pointers (`item_list` member variable) that is used to sort the objects in decreasing score order.

## 3.10 `Query` class

A `Query` represents a topic for which a set of voters or rankers submit their preference lists. For example, a `Query` may be a simple question like "*who is the best football player for 2022?*", or a more complex structure like a sequence of genes.

### Implementation files

The Query class is defined in the `src/Query.h` header file; its member functions are implemented in `src/Query.cpp`.

### Technical Details

The role of this object is central in the entire rank aggregation process, since it connects the input preference lists and the output aggregate list. In particular, a `Query` is connected to:

- an `Aggregator` that triggers the execution of a rank aggregation algorithm on the collection of the `InputLists`, and

- an `Evaluator` the evaluates the quality of the generated aggregate list with respect to a set of relevance judgments.

## 3.11 `Ranking` class

This is a simple class that stores the ranking information of a `MergedItem` in a particular input preference list. An array of `Ranking` objects is maintained inside each `MergedItem`.

### Implementation files

The Ranking class is defined in the `src/Ranking.h` header file; its member functions are implemented in `src/Ranking.cpp`.

### Technical Details

A Ranking class comprises three members:

- A pointer to the corresponding preference `InputList`,

- The (integer) ranking in the preference `InputList`, and

- The score value that is assigned by the voter to this element. The score data type can be either `float`, or `double`; this is determined by the `score_t` data type definition in `flagr.cpp`.

## 3.12 `Rel` class

A `Rel` object contains a relevance judgement about a list element. It is read from a special input CSV file and it is used by an `Evaluator` to compute several evaluation measures about the generated aggregate list.

A `Rel` object is a member of a `Rels` collection, which in turn is referenced by an `Evaluator`.

**Implementation files**

The `Rel` class is defined in the `src/Rel.h` header file; its member functions are implemented in `src/Rel.cpp`.

**Technical Details**

The `Rel` class consists of the following member variables:

- a string that represents the unique identifier of a `MergedItem`,

- an integer relevance judgment. The higher the value of this variable, the more relevant/important the item is considered.

- a next pointer to another `Rel` object that allows the creation of dynamic data structures (e.g. linked lists of `Rel` objects, etc.).

## 3.13 `Rels` class

`Rels` contains the relevance judgements that are required by the `Evaluator` in order to evaluate the quality of the generated aggregate list.

The relevance judgments are provided as an input to the library via a special CSV file.

**Implementation files**

The `Rels` class is defined in the `src/Rels.h` header file; its member functions are implemented in `src/Rels.cpp`.

**Technical Details**

Notice that the `Evaluator` contains a pointer to a `Rels` object. In this way, the `Evaluator` can quickly access the required relevance judgments during the evaluation procedure.

The `Rels` object is implemented as a standard hash table with the string item identifiers being its search key. The hash values are computed by the hash function of Daniel J. Bernstein, whereas the collisions are resolved by using chains in the form of linked lists.

## 3.14 `SimpleScoreStats` class

`SimpleScoreStats` is a very simple class that stores several score statistics. In the current FLAGR implementations, it is used exclusively for storing statistical information about the generated aggregate lists. More specifically, the minimum, maximum, and mean score values are stored there, including their standard deviation.

**Implementation files**

The SimpleScoreStats class is defined in the `src/SimpleScoreStats.h` header file; its member functions are implemented in `src/SimpleScoreStats.cpp`.

# 3.15 Voter class

A voter (also called ranker, or source), submits preferences for one or more topics (queries) in the form of a ranked preference list. The preference lists of all voters are subsequently aggregated by a rank aggregation method in order to generate a single consensus ranking. In FLAGR, a Voter submits a single `InputList` for each `Query`.

A Voter object possesses two properties:

- A unique string identifier that represents the voter's name, and

- A weight value that reflects the importance (degree of expertise) of the voter for a particular query. Non-weighted rank aggregation methods consider that all voters are equivalent. Therefore, their lists are processed in an identical manner. In contrast, the weighted methods apply unsupervised learning techniques and exploratory analysis to automatically determine the significance of each voter. The weight data type can be either `float`, or `double`; this is determined by the `score_t` data type definition in `flagr.cpp`.

**Implementation files**

The Voter class is defined in the `src/Voter.h` header file; its member functions are implemented in `src/Voter.cpp`.

# 3.16 Integrating custom methods

This step-by-step guide describes how custom rank aggregation methods can be implemented and integrated into FLAGR. If you intend to implement fewer than three methods, then several of the steps below are already implemented in FLAGR. For three or more custom methods, additional actions must be performed.

**Implementing your own method/s**

Custom methods must be implemented as C++ functions in the `src/ram/CustomMethods.cpp` file. This file already contains two such functions: CustomMethod1() and CustomMethod2(). In case you desire to implement more methods (e.g., CustomMethod3(), etc.), then you must implement them similarly in that file.

Typically, a rank aggregation method assigns scores to the list elements and then, it sorts the elements in either increasing, or decreasing score order. For this reason, the last step in your implementation must be the sorting of the elements of `MergedList`. Observe that CustomMethod1() and CustomMethod2() contain a call to qsort (QuickSort) in order to perform the required sorting.

Each algorithm implementation (including the built-in ones) takes three arguments:

- An array of the input preference lists (class InputList ** inlists). Most rank aggregation methods do not require access to this array, since when the function is called, the input lists have already been merged in the MergedList object. However, several methods require the computation of list distances (e.g. DIBRA, Agglomerative) and the inlists pointer provides access to this array.

- A pointer to a SimpleScoreStats object in case you desire to store score statistics (max, min, mean, etc.).

- A pointer to the InputParams object that contains the user-defined input parameters.

The following code example demonstrates an iteration through the elements of the aggregate list. For each element q, an iteration through its individual rankings in each input preference list is performed:

```cpp
void   MergedList::CustomMethod1   (class   InputList   **   inlists,   class
SimpleScoreStats * s, class InputParams * prms) {
      class MergedItem * q;
      class Ranking * r;

      for (rank_t i = 0; i < this->num_nodes; i++) {
              /// q stores an element of the aggregate list
              q = this->item_list[i];

              /// Iterate through the individual rankings of q
              for (uint32_t j = 0; j < q->get_num_alloc_rankings(); j++) {
                      r = q->get_ranking(j);
                      /// Do something with q and r

              }
      }
      /// Sort the list elements in decreasing score order
      qsort(this->item_list, this->num_nodes, sizeof(class MergedItem *),
              &MergedList::cmp_score_desc);
}
```

## Calling the new method/s

The implementation of the new methods can be immediately used, provided that you have not changed the names of the functions CustomMethod1() and CustomMethod2(). In this case, the following piece of code inside the main() function in cflagr.cpp and dllflagr.cpp invokes the new implementations:

```cpp
/// Execution of CustomMethod1
Custom1(input_file, qrels_file, 20, "RRA", output_dir);

/// Execution of CustomMethod2
Custom2(input_file, qrels_file, 20, "RRA", output_dir);
```

If you change the aforementioned function names, or you create a new function for a new algorithm implementation, then a sequence of actions must be taken so that it becomes available for usage. More specifically:

- The user must determine an integer identifier for the algorithm. To avoid conflicts with the built-in methods of FLAGR, it is advised that the leading digit of the identifier is 9 (e.g., 900).

- The user must update the aggregate() function of the Aggregator to allow the execution of the custom method. More specifically, the if statement must be appropriately extended.

- An exposed C function must be written in both cflagr.cpp and dllflagr.cpp with the aim of including the custom implementation in a shared/dynamic link library. The function can be called in the main function of cflagr.cpp.

## Additional details for custom method implementations

1. A custom rank aggregation method must be declared as a public member function of the MergedList class. This is performed in the class descriptor, in the src/MergedList.h header file. CustomMethod1() and CustomMethod2() are already public members of MergedList.

2. The required cpp file that contains the implementation of the member function must be stored in the src/ram directory.

3. Then, the cpp file is imported in the project with an include statement in src/MergedList.cpp.

# 4. PyFLAGR API Documentation

PyFLAGR is a Python library built on top of FLAGR. It includes a driver module called RAM that links to the FLAGR shared library and provides access to the algorithm implementations. Then, a set of classes inherit from RAM and allow the end user to execute the selected algorithm.



## 4.1 Introduction

PyFLAGR is a Python library built on top of FLAGR. It provides easy access to the algorithm implementations of FLAGR from standard Python programs. PyFLAGR has been designed with simplicity in mind: with only a few lines of code the programmer may efficiently execute complex rank aggregation methods and get the results in a Pandas Dataframe.

From a technical perspective, PyFLAGR links to the FLAGR shared/dynamic link library and makes use of its reference functions (namely, its exposed C functions) to pass the user-defined parameters and perform rank aggregation. In the sequel, it simply reads the output files that are produced by FLAGR and stores their contents in Pandas Dataframes. The Dataframes are then returned to the user.

PyFLAGR consists of the following modules:

- **RAM**: It implements the RAM base class that is responsible for several functional procedures like the library I/O, the linkage and loading of the FLAGR shared library, Dataframe handling, and so on. All the other classes of PyFLAGR derive from this class.

- **Linear**: It includes several classes that execute linear combination rank aggregation methods (CombSUM, CombMNZ, Borda Count).

- **Majoritarian**: It includes several classes that execute majoritarian rank aggregation methods (Condorcet Winners, Copeland Winners, Outranking Approach of Farah and Vanderpooten, 2007).

- **MarkovChains**: It executes the algorithms which are based on Markov Chains (MC1, MC2, MC3, MC4, see Dwork et al., 2001).

- **Weighted**: It executes the rank aggregation methods that automatically determine the voter weights in an unsupervised learning fashion (Preference Relations Method of Desarkar et al., 2016, Agglomerative Aggregation method of Chatterjee et al., 2018, Iterative Distance-based method of Akritidis et al., 2022).
- **Kemeny**: It executes Kemeny optimal aggregation (brute force, NP-Hard implementation).
- **RRA**: It executes the Robust Rank Aggregation (RRA) method of Kolde et al., 2012.

Please refer to the Publications section for more information about the relevant papers.

## 4.2 Installation

PyFLAGR can be installed directly by using pip:

```
pip install pyflagr
```

Alternatively, PyFLAGR can be installed from the sources by navigating to the directory where `setup.py` resides:

```
pip install /path/to/setup.py
```

## 4.3 `RAM` module

The RAM module implements a driver base class also named `RAM`. The `RAM` base class performs several important functional procedures, including PyFLAGR I/O, linkage and loading of the FLAGR shared library, I/O Dataframe handling, and so on. The majority of the other classes of PyFLAGR derive from this class and inherit its members and properties.

The class constructor takes as argument the `eval_pts` parameter that determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. The most important operation of the constructor is the loading of the FLAGR shared library, according to the underlying operating system. Therefore, if PyFLAGR is executed on a Linux-based system, then `pyflagr/pyflagr/flagr.so` is loaded. Similarly, if PyFLAGR is executed on a Windows-based system, then `pyflagr/pyflagr/flagr.dll` is loaded. For the time being, FLAGR has not been tested on MacOS-based systems and no pre-compiled shared libraries exist for this platform.

The successful loading of the FLAGR shared/dynamic link library creates the `flagr_lib` connection handler. `flagr_lib` acts as a connector between PyFLAGR and FLAGR, making the exposed C functions of FLAGR accessible from RAM and its derived classes.

Other member functions include `check_get_input()` and `check_get_rels_input()`. These two functions perform several sanity checks on the provided input files. On the other hand, the role of `get_output()` is to read the output files created by FLAGR and load their content into two Pandas Dataframes. These two Dataframes are eventually returned to the user.

**Implementation file**

`pyflagr/pyflagr/RAM.py`

# 4.4 `Linear` module

The `Linear` module provides access to the implementations of the linear combination methods of FLAGR. In these methods, the score of each element is computed by summing up the partial scores of that element with respect to its rankings in each input preference list. The module includes four classes which are described below: `CombSUM`, `CombMNZ`, `BordaCount` and `SimpleBordaCount`.

**Implementation file**

pyflagr/pyflagr/Linear.py

**The `CombSUM` and `CombMNZ` Python classes**

Both classes derive from RAM, a base class defined in the RAM module. They inherit the `flagr_lib` connector from RAM, and through it, they obtain access to the FLAGR shared library. Their constructors are identical and determine the data types of the input arguments and the return type of the `Linear()` exposed function. Observe the similarity between the members of `self.flagr_lib.Linear.argtypes` and the input arguments of the `Linear()` exposed function.

The arguments of the constructors of `CombSUM` and `CombMNZ` include:

| Parameter | Type | Default | Description |
|---|---|---|---|
| `eval_pts` | Integer, Optional. Considered only if `rels_file` or `rels_df` is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |
| `norm` | String, Optional. | borda | Rank or score normalization methods.<br><br>• `borda`: The aggregation is performed by normalizing the element rankings according to the Borda normalization method.<br>• `rank`: The aggregation is performed by normalizing the element rankings according to the Rank normalization method.<br>• `score`: The aggregation is performed by normalizing the element scores according to the Score normalization method.<br>• `z-score`: The aggregation is performed by normalizing the element scores according to the Z-Score normalization method.<br>• `simple-borda`: Similar to `borda` normalization but no partial score is assigned to an element if it is not ranked by a voter. |

`CombSUM` and `CombMNZ` also include an `aggregate()` function that receives the user-defined input parameters and passes them to the `Linear()` exposed C function, that subsequently performs the aggregation of the ranked input preference lists. The arguments of the `aggregate()` function include the following:

| Parameter | Type | Default | Description |
|---|---|---|---|
| input_file | String Required, unless input_df is set. | Empty String | A CSV file that contains the input lists to be aggregated. |
| input_df | Pandas DataFrame - Required, unless input_file is set. | None | A Pandas DataFrame that contains the input lists to be aggregated.<br><br>**Note:** If both input_file and input_df are set, only the former is used; the latter is ignored. |
| rels_file | String, Optional. | Empty String | A CSV file that contains the relevance judgements of the involved list elements. If such a file is passed, FLAGR will evaluate the generated aggregate list(s) by computing several retrieval effectiveness evaluation measures. The results of the evaluation will be stored in the eval_df DataFrame. Otherwise, no evaluation will take place and eval_df will be empty. |
| rels_df | Pandas DataFrame, Optional. | None | A Pandas DataFrame that contains the relevance judgements of the involved list elements. If such a dataframe is passed, FLAGR will evaluate the generated aggregate list(s) by computing several retrieval effectiveness evaluation measures. The results of the evaluation will be stored in the eval_df DataFrame. Otherwise, no evaluation will take place and eval_df will be empty.<br><br>**Note:** If both rels_file and rels_df are set, only the former is used; the latter is ignored. |
| output_dir | String, Optional. | Temporary directory (OS-specific) | The directory where the output files (aggregate lists and evaluation) will be stored. If it is not set, the default location will be used. |

### The `BordaCount` and `SimpleBordaCount` Python classes

These two classes have been included in PyFLAGR for historical reasons. They are equivalent to the CombSUM linear combination method with `norm='borda'` and `norm='simple-borda'` normalization methods, respectively.

BordaCount and SimpleBordaCount derive from CombSUM (which in turn derives from RAM). They do not have an `aggregate()` method, so they both call the same `aggregate()` function of CombSUM. Their only difference lies in their constructors: The former initializes CombSUM with `norm='borda'`, whereas the latter with `norm='simple-borda'`.

# 4.5 `Majoritarian` **module**

The Majoritarian module provides access to the implementations of the majoritarian rank aggregation methods of FLAGR. These methods are based on the majority criterion that, under several circumstances, identify the "winning" elements. The module includes three classes which are described below: `CondorcetWinners`, `CopelandWinners` and `OutrankingApproach`. Each method implements different scenarios for the majority criterion.

**Implementation file**

pyflagr/pyflagr/Majoritarian.py

**The `CondorcetWinners` and `CopelandWinners` Python classes**

Both classes derive from RAM, a base class defined in the RAM module. They inherit the `flagr_lib` connector from RAM, and through it, they obtain access to the respective exposed functions of FLAGR.

The constructor of `CondorcetWinners` determines the data types of the input arguments and the return type of the Condorcet() exposed function. Similarly, the constructor of `CopelandWinners` determines the data types of the input arguments and the return type of the Copeland() exposed function. Both constructors accept just one argument:

| Parameter | Type | Default | Description |
|-----------|------|---------|-------------|
| `eval_pts` | Integer, Optional. Considered only if `rels_file` or `rels_df` is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |

`CondorcetWinners` and `CopelandWinners` also include an `aggregate()` function that receives the user-defined parameters and passes them to the Condorcet() and Copeland() exposed functions respectively. The arguments of the `aggregate()` function are identical to those of the `CombSUM` and `CombMNZ` classes (refer to the respective table of Subsection 4.4).

**The `OutrankingApproach` Python class**

This class can be used to execute the Outranking Approach of Farah and Vanderpooten, 2007. Similarly to the other majoritarian methods, `OutrankingApproach` derives from RAM, a base class defined in the RAM module. It also inherits the `flagr_lib` connector from RAM and obtains access to the OutrankingApproach() exposed function of FLAGR.

The constructor of `OutrankingApproach` determines the data types of the input arguments and the return type of the OutrankingApproach() exposed function. Observe the similarity between the members of `self.flagr_lib.OutrankingApproach.argtypes` and the input arguments of the OutrankingApproach() exposed function.

| Parameter | Type | Default | Description |
|---|---|---|---|
| `eval_pts` | Integer, Optional. Considered only if `rels_file` or `rels_df` is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |
| `preference` | Hyper-parameter, Float, Optional. | 0.00 | The value of the preference threshold. |
| `veto` | Hyper-parameter, Float, Optional. | 0.75 | The value of the veto threshold. |
| `concordance` | Hyper-parameter, Float, Optional. | 0.00 | The value of the concordance threshold. |
| `discordance` | Hyper-parameter, Float, Optional. | 0.25 | The value of the discordance threshold. |

`OutrankingApproach` includes an `aggregate()` function that receives the user-defined parameters and passes them to the `OutrankingApproach() exposed function`. The arguments of the `aggregate()` function are identical to those of the `CombSUM` and `CombMNZ` classes (refer to the respective table of Subsection 4.4).

# 4.6 `MarkovChains` module

The `MarkovChains` module provides access to the implementations of the Markov Chains methods of Dwork et al., 2001 and DeConde et al., 2006. It includes the `MC` driver class, and five small derived classes (namely, `MC1`, `MC2`, `MC3`, `MC4` and `MCT`) that can be employed by the user to execute the corresponding methods.

**Implementation file**

`pyflagr/pyflagr/MarkovChains.py`

**The `MC` base class**

This is the driver class of the module. Its direct usage is weakly discouraged. The users should prefer employing the 5 classes that derive from `MC` (see below) and indirectly trigger the corresponding algorithm implementations of FLAGR.

Similarly to the other PyFLAGR classes, `MC` derives from `RAM`, another base class that is defined in the RAM module. It inherits the `flagr_lib` connector from RAM, and through it, it obtains access to the FLAGR shared library. Its constructor determines the data types of the input arguments and the return type of the FLAGR's `MC() exposed function`. Observe the similarity between the members of self.flagr_lib.MC.argtypes and the input arguments of the `MC() exposed function`.

The constructor of `MC` takes the following arguments:

| Parameter | Type | Default | Description |
|---|---|---|---|
| eval_pts | Integer, Optional. Considered only if rels_file or rels_df is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. For example, for eval_pts=10 FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |
| ergodic_number | Float, Optional. | 0.15 | The ergodic number, used during the computation of the ergodic transition matrix from the normalized transition matrix. |
| max_iterations | Integer, Optional. | 100 | The maximum number of iterations for the computation of the state matrix. |
| chain | Integer, Optional. | 804 | The Markov Chain method to execute. The possible values include:<br><br>• 801: Markov Chains Method 1 (MC1).<br><br>• 802: Markov Chains Method 2 (MC2).<br><br>• 803: Markov Chains Method 3 (MC3).<br><br>• 804: Markov Chains Method 4 (MC4).<br><br>• 805: Markov Chains Thurstone Method (MCT). |

MC includes an `aggregate()` function that receives the user-defined parameters and passes them to the MC() exposed function, that subsequently performs the aggregation of the ranked input preference lists. The arguments of the `aggregate()` function are identical to those of the `CombSUM` and `CombMNZ` classes (refer to the respective table of Subsection 4.4).

**The `MC1`, `MC2`, `MC3`, `MC4` and `MCT` derived classes**

These classes derive from the aforementioned `MC` base class; as children of `MC`, they also inherit from the RAM class. Each of these classes triggers the execution of a different Markov Chain algorithm, simply by passing different parameters to the constructor `MC`. Hence, `MC1` sets `chain=801`, `MC2` sets `chain=802`, and so on.

Also notice that notice that none of these classes have an `aggregate()` function. Consequently, the `aggregate()` of the base class (i.e. `MC`) is executed when the end user invokes that function.

# 4.7 `Weighted` module

The `Weighted` module provides access to the implementations of the weighted methods of FLAGR. Most rank aggregation approaches treat all input preference lists equally. In contrast, the weighted methods employ exploratory analysis techniques to automatically identify the expert voters in an unsupervised fashion. In the sequel, they assign higher weights to those who were identified as

experts, thus boosting the scores of their submitted elements. The module in question includes three classes: DIBRA, Agglomerative and PreferenceRelationsGraph.

## Implementation file

pyflagr/pyflagr/Weighted.py

## Distance-Based Iterative Rank Aggregation: The DIBRA Python class

This class links to the FLAGR implementation of the weighted method of Akritidis et al., 2022. It derives from RAM, a base class defined in the RAM module. It inherits the flagr_lib connector from RAM and through it, it obtains access to the FLAGR shared library. Its constructor determines the data types of the input arguments and the return type of the DIBRA() exposed function. Observe the similarity between the members of self.flagr_lib.DIBRA.argtypes and the input arguments of the DIBRA() exposed function.

The algorithm initially employs a standard non-weighted method to generate a starting consensus list. Then, it iteratively assigns converging weights to the voters according to the distances of their submitted lists with this consensus list. Therefore, the DIBRA constructor takes as arguments all the possible hyper-parameters of the supported non-weighted methods. Specifically:

| Parameter | Type | Default | Description |
|---|---|---|---|
| eval_pts | Integer, Optional. Considered only if rels_file or rels_df is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. For example, for eval_pts=10 FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |
| aggregator | Hyper-parameter, String, Optional. | combsum:borda | The selected non-weighted method that does the initial aggregation. Possible values include:<br>• combsum:borda: CombSUM with Borda normalization.<br>• combsum:rank: CombSUM with Rank normalization.<br>• combsum:score: CombSUM with minmax score normalization.<br>• combsum:z-score: CombSUM with Z-score normalization.<br>• combsum:simple-borda: CombSUM with simple Borda normalization.<br>• combmnz:borda: CombMNZ with Borda normalization.<br>• combmnz:rank: CombMNZ with Rank normalization.<br>• combmnz:score: CombMNZ with minmax score normalization.<br>• combmnz:z-score: CombMNZ with Z-score normalization. |

| | | | |
|---|---|---|---|
| | | | • `combmnz:simple-borda`: CombMNZ with simple Borda normalization.<br>• `condorcet`: Condorcet Winners.<br>• `copeland`: Copeland Winners.<br>• `outrank`: Outranking Approach. |
| `w_norm` | Hyper-parameter, String, Optional. | `minmax` | The voter weights normalization method. Possible values include:<br><br>• `none`: no weight normalization takes place.<br>• `minmax`: minmax weight normalization.<br>• `z`: z weight normalization. |
| `dist` | Hyper-parameter, String, Optional. | `cosine` | The correlation/distance metric that measures the distance between an input list and the temporary aggregate list. Possible values include:<br><br>• `rho`: Spearman's $\rho$.<br>• `cosine`: a metric based on cosine similarity (see Akritidis et al., 2022).<br>• `footrule`: Spearman's Footrule distance.<br>• `tau`: Kendall's $\tau$ |
| `prune` | Hyper-parameter, Boolean, Optional. | `False` | Triggers a weight-dependent list pruning mechanism. |
| `gamma` | Hyper-parameter, Float, Optional. | `1.5` | The $\gamma$ hyper-parameter that determines the steplength of weight learning. |
| `d1` | Hyper-parameter, Float, Optional. | `0.4` | The $\delta_1$ hyper-parameter of the list pruning mechanism. Applies only if `prune=True`. |
| `d2` | Hyper-parameter, Float, Optional. | `0.1` | The $\delta_2$ hyper-parameter of the list pruning mechanism. Applies only if `prune=True`. |
| `tol` | Hyper-parameter, Float, Optional. | `0.01` | Controls the convergence precision. This tolerance threshold represents the minimum precision of the difference between the voter weight in an iteration and the voter weight of the previous iteration. |
| `max_iter` | Hyper-parameter, Integer, Optional. | `50` | Controls the maximum number of iterations before the voter weights converge. |
| `pref` | Hyper-parameter, Float, Optional. | `0.0` | The preference threshold. Applies only if `aggregator=outrank`. |
| `veto` | Hyper-parameter, Float, Optional. | `0.75` | The veto threshold. Applies only if `aggregator=outrank`. |

| | | | |
|---|---|---|---|
| conc | Hyper-parameter, Float, Optional. | 0.0 | The concordance threshold. Applies only if aggregator=outrank. |
| disc | Hyper-parameter, Float, Optional. | 0.25 | The discordance threshold. Applies only if aggregator=outrank. |

DIBRA also includes an `aggregate()` function that receives the user-defined parameters and passes them to the DIBRA() exposed function, that subsequently performs the aggregation of the ranked input preference lists. The arguments of the `aggregate()` function are identical to those of the CombSUM and CombMNZ classes (refer to the respective table of Subsection 4.4).

## Agglomerative weighted aggregation: The `Agglomerative` Python class

This class employs the implementation of the Agglomerative weighted method of Chatterjee et al., 2018. Similarly to all weighted methods, it derives from RAM, a base class defined in the RAM module. It inherits the `flagr_lib` connector from RAM, and, through it, they obtain access to the FLAGR shared library. Its constructor determines the data types of the input arguments and the return type of the Agglomerative() exposed function. Observe the similarity between the members of `self.flagr_lib.Agglomerative.argtypes` and the input arguments of the Agglomerative() exposed function.

The constructor's arguments include:

| Parameter | Type | Default | Description |
|---|---|---|---|
| eval_pts | Integer, Optional. Considered only if `rels_file` or `rels_df` is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |
| c1 | Hyper-parameter, Float, Optional. | 0.1 | The $c_1$ hyper-parameter of the algorithm. |
| c2 | Hyper-parameter, Float, Optional. | 0.5 | The $c_2$ hyper-parameter of the algorithm. |

Agglomerative includes an `aggregate()` function that receives the user-defined parameters and passes them to the Agglomerative() exposed function, that subsequently performs the aggregation of the ranked input preference lists. The arguments of the `aggregate()` function are identical to those of the CombSUM and CombMNZ classes (refer to the respective table of Subsection 4.4).

## Preference relations weighted method: The `PreferenceRelationsGraph` Python class

This class links to the FLAGR implementation of the weighted method of Desarkar et al., 2016. Similarly to all weighted methods, it derives from RAM, a base class defined in the RAM module. It inherits the `flagr_lib` connector from RAM, and through it, they obtain access to the FLAGR shared library. Its constructor determines the data types of the input arguments and the return type of the PrefRel()

exposed function. Observe the similarity between the members of `self.flagr_lib.PrefRel.argtypes` and the input arguments of the `PrefRel()` exposed function. The constructor's arguments include:

| Parameter | Type | Default | Description |
|---|---|---|---|
| `eval_pts` | Integer, Optional. Considered only if `rels_file` or `rels_df` is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |
| `alpha` | Hyper-parameter, Float, Optional. | 0.1 | The $\alpha$ hyper-parameter of the algorithm. |
| `beta` | Hyper-parameter, Float, Optional. | 0.5 | The $\beta$ hyper-parameter of the algorithm. |

`PreferenceRelationsGraph` also includes an `aggregate()` function that receives the user-defined parameters and passes them to the `PrefRel()` exposed function, that subsequently performs the aggregation of the ranked input preference lists. The arguments of the `aggregate()` function are identical to those of the `CombSUM` and `CombMNZ` classes (refer to the respective table of Subsection 4.4).

# 4.8 Kemeny module

The Kemeny module provides access to the implementation of Kemeny Optimal Aggregation of FLAGR. It is stressed out that due to the method's high complexity, the brute force implementation becomes infeasible even when the number of elements, or the number of input lists receive moderate values.

**Implementation file**

`pyflagr/pyflagr/Kemeny.py`

**The `KemenyOptimal` class**

The `KemenyOptimal` class derives from RAM, a base class defined in the RAM module. It inherits the `flagr_lib` connector from RAM, and through it, it obtains access to the FLAGR shared library. Its constructor determines the data types of the input arguments and the return type of the FLAGR's `Kemeny()` exposed function. Observe the similarity between the members of `self.flagr_lib.Kemeny.argtypes` and the input arguments of the `Kemeny()` exposed function.

The arguments of the constructor includes one parameter:

| Parameter | Type | Default | Description |
|---|---|---|---|
| eval_pts | Integer, Optional. Considered only if `rels_file` or `rels_df` is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e., Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |

KemenyOptimal also includes an `aggregate()` function which receives the user-defined parameters and passes them to the Kemeny() exposed function, that subsequently performs the aggregation of the ranked input preference lists. The arguments of the `aggregate()` function are identical to those of the CombSUM and CombMNZ classes (refer to the respective table of Subsection 4.4).

## 4.9 RRA module

The RRA module provides access to the implementation of Robust Rank Aggregation of Kolde et al., 2012.

**Implementation file**

pyflagr/pyflagr/RRA.py

**The RRA class**

The RRA class derives from RAM, a base class defined in the RAM module. It inherits the `flagr_lib` connector from RAM, and through it, it obtains access to the FLAGR shared library. Its constructor determines the data types of the input arguments and the return type of the FLAGR's RobustRA() exposed function. Observe the similarity between the members of `self.flagr_lib.RobustRA.argtypes` and the input arguments of RobustRA() exposed function.

The arguments of the constructor includes two parameters:

| Parameter | Type | Default | Description |
|---|---|---|---|
| eval_pts | Integer, Optional. Considered only if `rels_file` or `rels_df` is set. | 10 | Determines the elements in the aggregate list on which the evaluation measures (i.e. Precision, and nDCG) will be computed. For example, for `eval_pts=10` FLAGR will compute Average Precision, $P@1, P@2, \ldots P@10$ and $N@1, N@2, \ldots N@10$. |
| exact | Boolean, Optional. | False | Determines whether the computed p-Values of the list elements will be corrected with the Stuart-Ares method. |

RRA also includes an `aggregate()` function that receives the user-defined parameters and passes them to the RobustRA() exposed function, that subsequently performs the aggregation of the ranked input preference lists. The arguments of the `aggregate()` function are identical to those of the CombSUM and CombMNZ classes (refer to the respective table of Subsection 4.4).

# 5. About

## 5.1 References

FLAGR utilizes algorithms, methods, and techniques from the following bibliography:

- E. Renda, U. Straccia, "Web metasearch: rank vs. score based rank aggregation methods", In Proceedings of the 2003 ACM symposium on Applied computing, pp. 841-846, 2003.

- M. Farah, D. Vanderpooten, "An outranking approach for rank aggregation in information retrieval", In Proceedings of the 30th ACM Conference on Research and Development in Information Retrieval, pp. 591-598, 2007.

- M.S. Desarkar, S. Sarkar, P. Mitra, "Preference relations based unsupervised rank aggregation for metasearch", Expert Systems with Applications, vol. 49, pp. 86-98, 2016.

- S. Chatterjee, A. Mukhopadhyay, M. Bhattacharyya, "A weighted rank aggregation approach towards crowd opinion analysis", Knowledge-Based Systems, vol. 149, pp. 47-60, 2018.

- L. Akritidis, A. Fevgas, P. Bozanis, Y. Manolopoulos, "An Unsupervised Distance-Based Model for Weighted Rank Aggregation with List Pruning", Expert Systems with Applications, vol. 202, pp. 117435, 2022.

- C. Dwork, R. Kumar, M. Naor, D. Sivakumar, "Rank Aggregation Methods for the Web", In Proceedings of the 10th International Conference on World Wide Web, pp. 613-622, 2001.

- R. Kolde, S. Laur, P. Adler, J. Vilo, "Robust rank aggregation for gene list integration and meta-analysis", Bioinformatics, vol. 28, no. 4, pp. 573-580, 2012.

- K.L. Majumder, G.P. Bhattacharjee, "Algorithm AS 63: The incomplete Beta Integral", Applied Statistics, vol. 22, no. 3, pp. 409-411, 1973.

- R.P. DeConde, S. Hawley, S. Falcon, N. Clegg, B. Knudsen, R. Etzioni, "Combining results of microarray experiments: a rank aggregation approach", Statistical Applications in Genetics and Molecular Biology, vol. 5, no. 1, 2006.

## 5.2 Acknowledgments

FLAGR, PyFLAGR and this Web site have been created by Leonidas Akritidis. You may contact me at lakritidis at ihu dot gr.

Deep thanks to John Burkardt for kindly permitting the usage of the ASA063 and ASA109 algorithm implementations about the computation of the incomplete beta integrals.

Bugs should be reported through the FLAGR's GitHub repository.

Third-party researchers are strongly encouraged to submit their own algorithm implementations, regardless of the programming language they are using.